

Computational Data Science

Joel Huang

September 25, 2018

1 Entity Relationship Diagrams & SQL

2 Lab: Data Handling in Unix

Previewing data

Raw data can be fairly large. We might want to examine the dataset's structure without opening it in a special program or text editor, which could take a long time and large amounts of memory.

Preview all - `cat`

Short for concatenate. Sequentially reads file(s) and writes them to `stdout`. If redirection `>` is used, then output is written to the specified file. `>` is used to write to a file and `>>` is used to append to a file.

```
cat file1.txt
cat file1.txt file2.txt > newcombinedfile.txt
cat >newfile.txt
cat -n file1.txt file2.txt > newnumberedfile.txt
cat file1.txt >> file2.txt
cat file1.txt file2.txt file3.txt | sort > test4
```

Preview some - `head`, `tail`

Use `head` or `tail` to preview the head or tail of the data. Remember to supply the flags:

```
-n Number of lines
-B Display number of lines before
-A Display number of lines after
```

Searching - `grep`

```
-E Use extended regular expression syntax
-o Output a matching segment of each line only
-n Print the line number of each matched line
-C Show a number of context lines too
```

```
. Matches any character.
* Matches zero or more instances of the preceding character.
+ Matches one or more instances of the preceding character.
[] Matches any of the characters within the brackets.
() Creates a sub-expression that can be combined to make more complicated expressions.
| OR operator; (www|ftp) matches either 'www' or 'ftp'.
^ Matches the beginning of a line.
$ Matches the end of the line.
\ Escapes the following character. Since . matches any character, to match a literal period you would need to use \.
```

3 Big Data, Hadoop, & MapReduce

4 Lab: MapReduce & Hadoop

Big Idea: Mappers & Reducers

5 Counting Relative Frequencies

Big Idea: Approximate probabilities by counting occurrences in the data

k-Nearest Neighbours

A way to measure similarity between different data points, through determining similarity values or distances. Classification with k-NN is straightforward, but sensitive to the value of k , potentially computationally expensive (but can be sped up using kd-tree), and takes memory to store all data points.

Decision trees & Random Forests

Each feature represents an opportunity for branching. Decision trees are inexpensive and explainable, but prone overfitting without pruning. To prevent overfitting, use Random forests-ensemble approach that aggregates decision trees's outputs via a majority voting system

Naive Bayes for classification

Assumption: Conditional independence

$$P(x_1, x_2, \dots, x_n | Y) = P(x_1 | Y) \cdot P(x_2 | Y) \cdot \dots \quad (1)$$

We are trying to find the most likely class given an observation Maximum Most likely class Y given feature set X

$$y_{MAP} = \arg \max P(Y | x_1, x_2, \dots, x_m) \quad (2)$$

$$= \arg \max \frac{P(X | Y) P(Y)}{P(X)} \quad (3)$$

In equation (2), we're comparing the arg max of $P(Y = 0 | X)$ and $P(Y = 1 | X)$. Therefore we can cancel the constant denominator $P(X)$:

$$= \arg \max P(X | Y) P(Y) \quad (4)$$

Next, we need to find $P(X | Y)$ and $P(Y)$.

Count the prior probability $P(Y)$ Class labels are $Y \in \{0, 1\}$. Therefore, for a dataset with n labels,

$$P(Y = 0) = \frac{n_{Y=0}}{n} \quad (5)$$

$$P(Y = 1) = \frac{n_{Y=1}}{n} \quad (6)$$

In other words, there are $n_{Y=1}$ out of n occurrences of label Y having value 1.

Estimate $P(X | Y = 0)$ and $P(X | Y = 1)$ Since we have assumed conditional independence of X given classes Y (1),

$$P(x_i | Y = 0) = \frac{|x_{i,Y=0}|}{n_{Y=0}} \quad (7)$$

We are counting the probability (occurrences in the dataset) of feature x_i given $Y = 0$. Find all the rows where $Y = 0$ and count the occurrences of feature x_i . Repeat this for $Y = 1$.

Final comparison (arg max) Now we have pairs $P(Y = 0)$, $P(X | Y = 0)$, and $P(Y = 1)$, $P(X | Y = 1)$. Multiply the pairs as in equation (4) and take the *max* of the two.

Naive Bayes for word counting

Assumptions: Conditional independence, irrelevance of word order

We want to predict a class for a given sentence, for example, "good" for the sentence "love the burgers here, delicious and filling". We can estimate the class c using Naive Bayes,

$$c = \arg \max_{c_i \in C} P(c_i) \prod_{w_j \in W} P(w_j | c_i) \quad (8)$$

For a set of documents D , words W , and classes C , the probability of class c_i is simply the number of occurrences of c_i in document D .

$$P(c_i) = \frac{|c_i|}{|D|} \quad (9)$$

The probability of word w_1 given class c_i is also simply the number of occurrences of word w_1 in sentences with class c_i , divided by the probability of class c_i .

$$\begin{aligned} P(w_1 | c_i) &= \frac{P(w_1, c_i)}{P(c_i)} \\ &= \frac{\text{count}(w_1, c_i)}{\sum_{w_j \in W} \text{count}(w_j, c_i)} \end{aligned} \quad (10)$$

Example: Restaurant ratings

ID	Sentence	Class
1	The burger is tasteless and slow service	Bad
2	slow serving time and everything is horrible	Bad
3	Restaurant is near MRT, serves delicious burgers	Good
4	love the burger here, delicious and filling	Good
5	love this place, delicious burgers but slow service	?

Where classes $C = \{Bad, Good\}$, and words $W = \{tasteless, slow, horrible, delicious, love\}$.

$$P(c) = \begin{cases} \frac{2}{4}, & c = Bad \\ \frac{2}{4}, & c = Good \end{cases} \quad (11)$$

$$\begin{aligned} P(love | Good) &= \frac{\text{count}(love, Good)}{\sum_{w_j \in W} \text{count}(w_j, Good)} = \frac{1}{3} \\ P(delicious | Good) &= \frac{\text{count}(delicious, Good)}{\sum_{w_j \in W} \text{count}(w_j, Good)} = \frac{2}{3} \end{aligned} \quad (12)$$

*** TODO: CHECK THIS**

$$\begin{aligned}
P(\text{Good} \mid \text{love}, \text{delicious}) &= P(\text{Good}) \prod_{w_j \in W} P(w_j \mid \text{Good}) \\
&= P(\text{Good}) \cdot P(\text{love} \mid \text{Good}) \cdot P(\text{delicious} \mid \text{Good}) \\
&= \frac{2}{4} \times \frac{1}{3} \times \frac{2}{3} = \frac{1}{9}
\end{aligned}
\tag{13}$$

$$\begin{aligned}
P(\text{tasteless} \mid \text{Bad}) &= \frac{\text{count}(\text{tasteless}, \text{Bad})}{\sum_{w_j \in W} \text{count}(w_j, \text{Bad})} = \frac{1}{4} \\
P(\text{slow} \mid \text{Bad}) &= \frac{\text{count}(\text{slow}, \text{Bad})}{\sum_{w_j \in W} \text{count}(w_j, \text{Bad})} = \frac{2}{4} \\
P(\text{horrible} \mid \text{Bad}) &= \frac{\text{count}(\text{horrible}, \text{Bad})}{\sum_{w_j \in W} \text{count}(w_j, \text{Bad})} = \frac{1}{4}
\end{aligned}
\tag{14}$$

*** TODO: INCOMPLETE**

$$\begin{aligned}
P(\text{Bad} \mid \text{love}, \text{delicious}) &= P(\text{Bad}) \prod_{w_j \in W} P(w_j \mid \text{Bad}) \\
&= P(\text{Bad}) \cdot P(\text{love} \mid \text{Bad}) \cdot P(\text{delicious} \mid \text{Bad}) \\
&= \frac{2}{4} \times \frac{1}{3} \times \frac{2}{3} = \frac{1}{9}
\end{aligned}
\tag{15}$$